

1. a) For nested-loop, the biggest advantage is that it requires no indices and can be used with any kind of join condition. However, it is expensive as it will examine every pair of tuples in  $r \bowtie s$ , and for the worst case where memory is only enough for holding one block of each relation, cost will be  $(n_r * b_s + b_r) + (n_r + b_r)$

For merge-join, it is efficient when dealing with the situation where both relations are already sorted on the join key, and for sequential disk access pattern, it can be faster and reduce seek time. However, if the input relations are not sorted, the cost of sorting can be significant especially for large relations where sorting might involve additional disk I/O

b) For  $r$

① run 1

load 3 entries in 3 blocks in buffer

$(77, c)$ ,  $(77, a)$ ,  $(77, g)$

sort based on B

$(77, a)$ ,  $(77, c)$ ,  $(77, g) \Rightarrow$  run 1

② run 2

load 3 entries in 3 blocks in buffer

$(31, h)$ ,  $(40, f)$ ,  $(3, d)$

sort based on B

$(3, d), (40, f), (31, h) \Rightarrow \text{run 2}$

③ merge  $N=2 < M=3$ , two blocks to buffer inputs, 1 blocks to buffer output

input	input	Out Put
$(77, a)$	$(3, d)$	$(77, a)$
$(77, c)$	$(3, d)$	$(77, c)$
$(77, g)$	$(3, d)$	$(3, d)$
$(77, g)$	$(40, f)$	$(40, f)$
$(77, g)$	$(31, h)$	$(77, g)$
		$(31, h)$

For S

For each run: Pick 3 entries into 3 blocks of buffer, sort on B

① run 1

$(a, 2), (d, 1), (g, 5)$

② run 2

$(b, 6), (e, 10), (f, 2)$

③ run 3

$(c, 9), (h, 7)$

④ Merge  $N = M = 3$

1. merge  $M-1 = 2$  runs

take run 1 & run 2

input	input	output
(a, 2)	(b, 6)	(a, 2)
(d, 1)	(b, 6)	(b, 6)
(d, 1)	(e, 10)	(d, 1)
(g, 5)	(e, 10)	(e, 10)
(g, 5)	(f, 2)	(f, 2)
		(g, 5)
		run 4

⑤ merge  $N=2 < M$

take run 4 & run 3

input	input	output
(a, 2)	(c, 9)	(a, 2)
(b, 6)	(c, 9)	(b, 6)
(d, 1)	(c, 9)	(c, 9)
(d, 1)	(h, 7)	(d, 1)
(e, 10)	(h, 7)	(e, 10)
(f, 2)	(h, 7)	(f, 2)
(g, 5)	(h, 7)	(g, 5)
		(h, 7)

$$\therefore$$

r	A	B
	77	a
	77	c
	3	d
	40	f
	77	g
	31	h

s	B	C
	a	2
	b	6
	c	9
	d	1
	e	10
	f	2
	g	5
	h	7

For merge-join

1<sup>st</sup> I/O: load (a, 2)

2<sup>nd</sup> I/O: load (77, a)

4<sup>th</sup> I/O: load (b, b)

5<sup>th</sup> I/O: load (77, c)

6<sup>th</sup> I/O: load (c, 9)

8<sup>th</sup> I/O: load (d, 1)

9<sup>th</sup> I/O: load (3, d)

11<sup>th</sup> I/O: load (e, 10)

12<sup>th</sup> I/O: load (40, f)

13<sup>th</sup> I/O: load (f, 2)

15<sup>th</sup> I/O: load (g, 5)

16<sup>th</sup> I/O: load (77, g)

$\Rightarrow$  3<sup>rd</sup> I/O: write (a, 2, 77)

$\Rightarrow$  7<sup>th</sup> I/O: write (c, 9, 77)

$\Rightarrow$  10<sup>th</sup> I/O: write (d, 1, 3)

$\Rightarrow$  14<sup>th</sup> I/O: write (f, 2, 40)

$\Rightarrow$  17<sup>th</sup> I/O write (g, 5, 77)

18<sup>th</sup> I/O: load (h, 7)  $\Rightarrow$  20<sup>th</sup> I/O write (h, 7, 31)  
 19<sup>th</sup> I/O: load (31, h)

2.  $M = 3$ , each 2 records

load 6 records on 3 blocks

10, 11, 1, 5, 90, 1

sort

1, 1, 5, 10, 11, 90  $\Rightarrow$  run 1

load 2 records on 1 block

2, 101  $\Rightarrow$  run 2

a) 2 runs

run 1: (1, 1), 5, 10, 11, 90)

run 2: (2, 101)

b)  $\therefore N = 2$ ,  $M = 3$   $N < M$

$\therefore$  3 blocks, 1 for run 1 input, 1 for run 2 input,  
 1 for output

input	input	output
(1, 1) $\leftarrow$ load	(2, 101) $\leftarrow$ load	(1, 1)
(5, 10) $\leftarrow$ load	(2, 101)	(2, 5)
10	101	10
(11, 90) $\leftarrow$ load	101	(11, 90)

10]

10]

Output: 1, 1, 2, 5, 10, 11, 90, 101

Using N-way merge

C) for creating runs

load  $3 + 1 = 4$  blocks from disk to memory

for merging

run 2 fit in 1 block  $\Rightarrow$  1 block transferred

run 1 loads 3 blocks  $\Rightarrow$  3 blocks transferred

$$\therefore 3 + 1 = 4$$

$$\therefore \text{total} = \text{merge} + \text{create} = 4 + 4 = 8 \text{ blocks}$$