

1. First it will scan from bottom to top to find the most recent checkpoint, which is checkpoint (T_0)

We have transaction T_0 and T_1


T_0 begins before checkpoint, but never committed or aborted \Rightarrow undo

T_1 begins after checkpoint and committed after checkpoint \Rightarrow redo

To undo T_0 : $\langle T_0, B, 2000 \rangle \langle T_0, A, 1000 \rangle \langle T_0, \text{abort} \rangle$
will be written

To redo T_1 : Z will be set to 150, but no log will be written

2. For testing for conflict serializability



$T_1 \rightarrow T_2$: as T_1 read(A) T_2 write(A) \therefore conflict

$T_2 \rightarrow T_3$: as T_2 read(B) T_3 write(B) \therefore conflict

$T_3 \rightarrow T_1$: as T_3 read(A) T_1 write(A) \therefore conflict

$T_3 \rightarrow T_2$: as T_3 read(A) T_2 write(A) \therefore conflict

It is not a cyclic

\therefore not conflict serializable

3. a) Since in T_1 it will release the X-lock right after $\text{write}(A)$. Then T_2 will have a chance to hold a S-lock on A , and to read the updated A by T_1 . T_2 might potentially applied A to write B which then will be read by T_3 . In this case, since update on A in T_1 is not committed, if T_1 abort, A will be rolled back to the state before any updates happens. However, T_2 & T_3 are already using updated version of A . Thus, it will result a cascading rollback of T_2 and T_3 as well, as they are using an invalid A for potential write.

b) strict 2PL meaning T_1 will not release X-lock on A until abort. So that T_2 which acquires S-lock on A has to wait for T_1 's X-lock to be released and T_3 will wait for T_2 . In this case, only if T_1 is committed which means the update is valid or aborted, T_2 and T_3 can have a chance to read (A) and make modification further. Thus, T_2 and T_3 will never have a chance to use an invalid A to make updates. So, they don't need to be rolled back.

4. Cascading abort occur when one transaction ^(T₁) aborts after making some changes that other transactions ^(T₂ & T₃) have already read. In this case If T₁ aborts, any other transactions that has read its uncommitted changes must also abort to maintain consistency.

But for the case where all X-locks are held until after the transaction holding the lock committed or aborted, and while holding X-locks, S-locks and further X-locks are not allowed to be acquired on the same items. In this case, other transactions are prevented from accessing the affected/updated items before it is committed. This eliminate the possibility of reading uncommitted data to improve Isolation, atomicity and Consistency which eliminates the needs & risk of cascading aborts.